

Section 1: The 80/20 Core (Hiring-Critical)

1. Authentication ≠ Authorization: The Boundary That Kills Production Systems

Why hiring managers care: Most IAM incidents trace to confused responsibility between "who you are" (authn) vs "what you can do" (authz). Engineers who conflate these create brittle systems that break during federation, fail during privilege changes, and create security gaps.

How platforms implement it:

- **Okta/Azure AD:** Authentication produces tokens (ID token for identity claims, access token for API authorization). The *identity provider* doesn't decide if you can delete an S3 bucket—it asserts who you are.
- **AWS IAM:** IAM policies (authorization) assume identity is already proven (by federation, IAM user, or role assumption). The STS service bridges authentication → authorization.
- **Ping/ForgeRock:** PingFederate does SAML/OIDC authentication; PingAccess does authorization enforcement. Separate deployment topologies, separate failure modes.

What breaks when you misunderstand:

- **Broken SSO scenario:** Team puts business logic ("can this user access beta features?") in the IdP. Now every app must call back to the IdP for every permission check. IdP becomes a single point of failure for *authorization*, not just authentication. Under load, entire platform fails.
- **Privilege escalation:** App trusts `groups` claim in ID token without verifying token signature or checking revocation. Attacker modifies JWT locally, gains admin access.
- **Operational fragility:** Authorization rules scattered across IdP, app code, and API gateways. Nobody knows the complete picture. Compliance audit fails.

The hiring signal: Can you draw a sequence diagram showing where authentication ends and authorization begins in a multi-tier app using OAuth2 + OIDC? Can you explain why access tokens should be opaque to clients but ID tokens must be JWTs?

2. Token Lifecycle Reality: The Phantom Menace of "Instant Revocation"

Why hiring managers care: "We need to revoke access immediately when someone is terminated" is the #1 business requirement that collides with distributed systems reality. Engineers who promise instant revocation without understanding token mechanics create false security.

How platforms handle it:

- **Short-lived access tokens (5-60 min):** Okta, Azure AD, Keycloak default to this. Revocation is "eventually consistent" up to token TTL.
- **Refresh token rotation:** Okta/Azure AD rotate refresh tokens on use. Stolen refresh tokens have limited blast radius *if* rotation is enforced.
- **Reference tokens vs JWT:** Ping/ForgeRock let you choose. JWT = no revocation check (faster, scales). Opaque token = must call back to IdP (slower, revocable). Most enterprises use JWT and accept delayed revocation.
- **Session management layer:** Azure AD session cookies (hours/days) *above* access tokens (minutes). Revoking the session doesn't kill the access token.

What breaks:

- **The M&A disaster:** Company acquires another company. Need to onboard 5,000 users by Monday. Choose 4-hour access tokens for "security." Users can't stay logged in, flood helpdesk. Nobody understood refresh token flows.
- **The terminated employee:** Fire someone at 9am. Their access token is valid until 10am (1-hour TTL). They exfiltrate data at 9:45am. Business blames "IAM didn't revoke immediately."
- **Cascading logout failure:** App only clears local session cookie. Doesn't call IdP logout endpoint. User "logs out" but can still access federated apps. Nobody implemented back-channel logout (OIDC) or single logout (SAML).

The hiring signal: Can you explain why continuous access evaluation (Azure AD) or token binding (Okta) exist? Can you design a system that achieves "revocation within 5 minutes" without calling the IdP on every API request?

3. Identity as a Distributed System: CAP Theorem Applies

Why hiring managers care: Identity systems span multiple data centers, clouds, and SaaS apps. Engineers who treat identity as a monolithic database design systems that split-brain during network partitions or fail to scale.

How platforms handle it:

- **Okta:** Multi-region active-active, but *not* multi-master. Primary region for writes, read replicas elsewhere. During region failure, authentication can succeed (cached), but provisioning/de-provisioning blocks.
- **Azure AD:** Global directory, replicated across regions. Writes go to master, propagate to replicas (eventual consistency ~seconds to minutes). Group membership changes can lag.
- **Ping Directory:** LDAP replication with configurable conflict resolution. Multi-master possible but operationally complex. Most enterprises run active-passive for writes.
- **AWS IAM:** Eventually consistent. IAM policy change in us-east-1 takes seconds to propagate globally. Security teams often don't know this.

What breaks:

- **The split-brain scenario:** Enterprise runs Ping on-prem (primary) and Okta cloud (disaster recovery). Network partition. Both accept authentications. User logs into on-prem (success), then cloud app (fails because group membership not synced). Business thinks SSO is "broken."
- **The race condition:** Provision user → assign group → user tries to login. In Azure AD, user exists but group membership not propagated (eventual consistency). User gets 403. Retry after 30 seconds works. Engineering adds "sleep(5)" hacks everywhere.
- **The federation loop:** User authenticates with Okta → redirects to Azure AD (federation) → redirects back to Okta (nested federation). Each IdP has caching behavior. User sees different attributes at each layer. Debugging impossible without distributed tracing.

The hiring signal: Can you explain how Okta's org2org replication differs from Azure AD Connect? When would you choose synchronous LDAP replication vs async REST API syncing?

4. Federation Trust Modeling: Blast Radius of Misconfigured Trust

Why hiring managers care: Federation is how enterprises scale identity (partners, customers, acquired companies). Misconfigured trust relationships create backdoors. Most breaches involving SAML/OIDC trace to bad trust boundaries.

How platforms implement it:

- **SAML:** Trust is symmetric and explicit. SP trusts IdP via metadata exchange (certificates, endpoints). Each relationship is manually configured. Scales poorly (N×M problem).
- **OIDC:** Discovery via `.well-known/openid-configuration`. Trust is asymmetric (SP trusts IdP public keys). Scales better but requires PKI discipline.
- **Azure AD multi-tenant apps:** Service principal in each tenant. Misconfigure consent scope, app gains god-mode in all tenants.
- **AWS IAM cross-account:** Trust policy in target account + assume role in source account. Misconfigure, break principle of least privilege.

What breaks:

- **The SolarWinds-style attack:** Attacker compromises IdP signing certificate. Can now issue valid SAML assertions for *any* user to *any* SP that trusts that IdP. Blast radius: every federated app.
- **The overprivileged app:** Azure AD app requests `Directory.ReadWrite.All` scope. Tenant admin consents. App can now modify groups, users, conditional access policies. Intended to only read user profiles.
- **The transitive trust nightmare:** Company A trusts Company B (partner). Company B trusts Company C (vendor). User from Company C can now access Company A resources. Nobody mapped the transitive trust.

The hiring signal: Can you design a federation architecture for M&A where acquired company IdP is distrusted after 90 days? How do you prevent one compromised federated partner from affecting others?

5. Group and Role Explosion: Entropy of Entitlements

Why hiring managers care: Every enterprise IAM system starts clean and devolves into thousands of groups with unclear purpose. This creates security risk (privilege creep), operational cost (manual reviews), and audit failures.

How platforms handle it:

- **Azure AD:** Groups (security principals) vs administrative units (RBAC scoping). Role-assignable groups (limited to 500/tenant) vs regular groups (unlimited). Teams implicitly creates groups.
- **Okta:** App-specific groups vs universal directory groups. Group push to apps (provisioning) vs group-based assignment (authorization). No automatic cleanup.

- **AWS IAM:** Managed policies (AWS-owned) vs customer-managed vs inline. 10 policies/user or role limit forces composition. Group nesting not supported.
- **Ping Directory:** LDAP groups can nest infinitely. Resolving nested group membership becomes $O(n^2)$ problem at scale.

What breaks:

- **The 10,000 group problem:** Company has 10k groups in Azure AD. Half have no owner. Conditional access policy uses group membership. Can't delete groups because "something might break." Attack surface: massive.
- **The privilege creep:** User joins Team A (gets group A). Moves to Team B (gets group B, keeps A). Repeats 5x over 3 years. Now has 5 groups, each with admin rights to different systems. Nobody knows their effective permissions.
- **The orphaned entitlement:** Okta pushes groups to Salesforce via SCIM. Group deleted in Okta. SCIM only provisions, doesn't de-provision group membership. User retains access. Audit fails.

The hiring signal: Can you design a group naming taxonomy that prevents this? How do you implement periodic attestation without manual spreadsheets?

6. Human vs Workload Identity: Divergent Evolution

Why hiring managers care: Same "identity" word, completely different requirements. Human identity: MFA, sessions, interactive. Workload identity: API keys, short-lived tokens, no interactivity. Mixing them creates security and operational problems.

How platforms handle it:

- **Azure AD:** Managed identities (VM, Function) get Azure AD tokens without secrets. Service principals (app registrations) use client secrets or certificates. Different RBAC model.
- **AWS IAM:** IAM users (humans, long-term creds) vs IAM roles (workloads, temporary creds via STS). Roles don't have passwords/MFA.
- **Okta:** Service apps use OAuth2 client credentials flow. No user context, no MFA. Separate token endpoint.
- **Keycloak:** Service accounts feature for client credentials. Separate from user realms.

What breaks:

- **The shared service account:** Team creates "[svc-app1@company.com](#)" in Azure AD, shares password in wiki. App uses it for API calls. Account has no MFA (can't do non-

interactive MFA). Account has `User.ReadWrite.All` because "the app needs to create users." Security nightmare.

- **The long-lived token:** CI/CD pipeline stores Okta API token (no expiration) in GitHub repo. Developer forks repo. Token leaked. Attacker has permanent admin access to Okta.
- **The workload MFA disaster:** Require MFA for all Azure AD sign-ins. Break every CI/CD pipeline because service principals can't do MFA. Emergency rollback at 2am.

The hiring signal: Can you design a migration from service accounts to managed identities? How do you handle secrets rotation for legacy apps that can't use OIDC?

7. Just-In-Time vs Pre-Provisioning: The Cold Start Problem

Why hiring managers care: JIT provisioning sounds great ("no pre-provisioning!") until you hit edge cases. Pre-provisioning scales but creates drift. Most enterprises run hybrid and suffer both problems.

How platforms handle it:

- **Okta:** JIT via SAML attribute mapping or OIDC claims. User doesn't exist → create on first login. Works until app requires pre-existing user ID in foreign key relationships.
- **Azure AD:** B2B guests are JIT. SCIM to SaaS is pre-provisioning. Hybrid creates races.
- **SCIM standard:** Idempotent provisioning, but apps implement it differently. Some apps JIT on top of SCIM (chaos).

What breaks:

- **The referential integrity problem:** Salesforce contact records reference user IDs. JIT user on first login → user doesn't exist when contact created → foreign key violation. Pre-provisioning required.
- **The offboarding gap:** JIT creates users but never deletes them. User leaves company, Okta account disabled. User still exists in 50 SaaS apps (just not able to login via SSO). Direct password reset still works.
- **The cold start penalty:** 10,000 users need access to new app. Choose JIT. First 10,000 logins take 5 seconds each (account creation). Users think app is slow. Should have pre-provisioned.

The hiring signal: When do you choose JIT vs SCIM provisioning? How do you handle apps that require pre-existing users for licensing?

8. Session Management Reality: The Illusion of Single Logout

Why hiring managers care: Users expect "logout" to mean "logged out everywhere." Reality: sessions span multiple layers (browser cookie, IdP session, app session, token lifetime). True single logout is hard.

How platforms handle it:

- **SAML:** Single Logout (SLO) requires every SP to register logout endpoints. IdP sends logout requests to all SPs in parallel. If one SP is down, logout fails silently.
- **OIDC:** Back-channel logout (HTTP POST to apps) or front-channel logout (iframes). Both require app cooperation. Most apps don't implement.
- **Azure AD:** Session lifetime separate from token lifetime. Revoking session doesn't invalidate existing access tokens.

What breaks:

- **The phantom session:** User clicks logout in App A. App clears local cookie, doesn't notify IdP. User still logged into IdP, can access App B without re-auth. User thinks they logged out.
- **The iframe nightmare:** Enterprise implements front-channel logout (OIDC). Requires loading iframe per app. User has 20 apps. Logout page loads 20 iframes. Browser blocks some as cross-origin. Logout fails silently for those apps.
- **The mobile trap:** Mobile app stores refresh token in secure enclave. User logs out → app deletes token. But server-side session still valid. Token replay attack possible if stolen.

The hiring signal: Can you design a logout flow that achieves "reasonable" user expectation without perfect SLO? How do you handle mobile apps in logout scenarios?

9. IdP Chaining and Delegation: Amplification of Complexity

Why hiring managers care: Enterprises rarely have one IdP. Okta fronts Azure AD fronts on-prem AD. Each layer adds latency, failure modes, and attribution complexity. Most outages involve chaining.

How platforms handle it:

- **Okta → Azure AD federation:** Okta acts as SP to Azure AD (IdP). Okta issues its own token after Azure auth. Attribute mapping at both layers.
- **Azure AD → ADFS → AD:** Azure AD Connect syncs AD → Azure AD. ADFS federates to Azure AD for specific apps. Three sources of truth.

- **Ping → LDAP:** PingFederate uses PingDirectory (LDAP) as user store. Separate authentication (Ping) from directory (LDAP). Schema mapping required.

What breaks:

- **The attribution nightmare:** User can't login. Where did it fail? On-prem AD? Azure AD Connect sync? Azure AD conditional access? Okta federation? App-level authorization? Each layer logs differently.
- **The latency cascade:** User auth requires: DNS → Okta (50ms) → Azure AD federation (200ms) → ADFS (150ms) → on-prem AD (100ms) → back up the chain. Total: 500ms + network variance. Timeout set to 400ms. Random auth failures.
- **The infinite loop:** Okta trusts Azure AD. Azure AD trusts Okta (different apps). Misconfigured, user enters IdP loop: Okta → Azure → Okta → Azure. Browser eventually times out.

The hiring signal: Can you debug a failed SAML assertion that traversed 3 IdPs? What tools do you need?

10. Audit and Compliance as Operational Concerns, Not Checkboxes

Why hiring managers care: "We need audit logs for SOC2" is the *start*, not the end. Real compliance means being able to answer "who accessed what, when, and why" in under 5 minutes during an incident. Most enterprises can't.

How platforms handle it:

- **Azure AD sign-in logs:** 30-day retention default. Exporting to Log Analytics costs money. Queries are slow (Kusto). No cross-cloud correlation by default.
- **Okta system logs:** Comprehensive, but rate-limited API for extraction. Third-party SIEM required for long-term storage.
- **AWS CloudTrail:** IAM actions logged, but must be explicitly enabled for data events. Cross-account aggregation requires setup.

What breaks:

- **The incident response failure:** Suspicious login from Russia at 3am. Need to find: what apps they accessed, what data they downloaded, which API keys they created. Azure AD logs show successful auth. App logs show 403 errors. AWS CloudTrail shows S3 access. No correlation. Incident response takes 8 hours instead of 5 minutes.
- **The compliance gap:** Auditor asks "show me all admin role changes in last 12 months." IAM team provides Azure AD logs (role assignments). Missed: AWS IAM policy changes

(different log), Okta super admin actions (different platform), app-level admin grants (no logs). Fail audit.

- **The privacy violation:** GDPR delete request for user. Delete from Azure AD, Okta, AWS IAM. Forget to delete audit logs. Logs still contain PII. GDPR violation.

The hiring signal: Can you design a unified audit architecture that spans Azure AD + AWS + Okta? What's your query response time SLA for "all actions by user X"?